

Project_Netflix

I will first begin by splitting the Netflix data into the training and test data sets.

```
library(ISLR2)
library(tree)
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.4.0      v purrr  1.0.1
## v tibble  3.1.8      v dplyr  1.1.0
## v tidyr   1.3.0      v stringr 1.5.0
## v readr   2.1.3      v forcats 1.0.0
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(caret)

## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##   lift

Netfl <- read.csv("Best Movies Netflix.csv")

Net <- subset(Netfl, select = c(RELEASE_YEAR:MAIN_GENRE))

Net$MAIN_GENRE <- as.factor(Net$MAIN_GENRE)

Netflix<-Net%>%
  as_tibble()

set.seed(456)

netflix_index = sample(1:nrow(Netflix), nrow(Netflix)/2)

NetflixTrain_set = Netflix[netflix_index,]

NetflixTest_set = Netflix[-netflix_index,]
```

The regression tree will only work on factor variables and numeric variables, so the 'MAIN_GENRE' variable was changed in order to prevent NA's introduced by coercion. This is also why the 'TITLE' variable and

'MAIN_PRODUCTION' variable were not included because transforming these variables created too many factors.

The regression tree will be fitted.

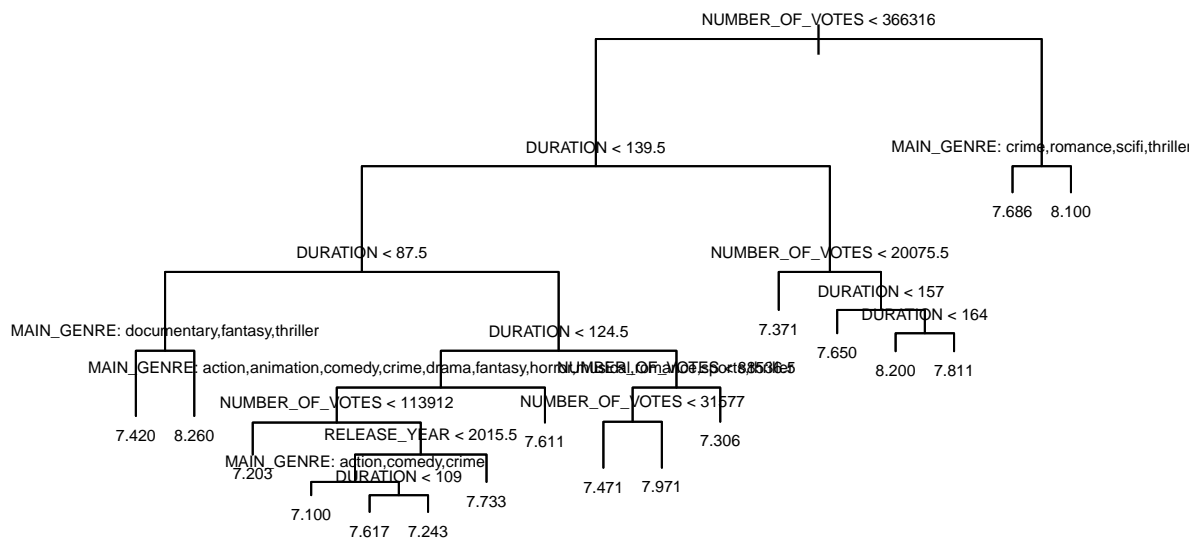
```
Netflix_regressiontree <- tree(SCORE~., NetflixTrain_set)
```

```
summary(Netflix_regressiontree)
```

```
##
## Regression tree:
## tree(formula = SCORE ~ ., data = NetflixTrain_set)
## Number of terminal nodes: 17
## Residual mean deviance: 0.08655 = 15.23 / 176
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -0.61110 -0.21110 -0.00625  0.00000  0.20000  0.75000
```

This tree was plotted in order to develop a visualization of it.

```
plot(Netflix_regressiontree)
text(Netflix_regressiontree,pretty=0,cex =0.5)
```



This regression tree will also be pruned with cross-validation.

```

set.seed(456)

Pruned_Netflix <- cv.tree(Netflix_regressiontree)

Pruned_Netflix

## $size
## [1] 17 15 14 13 12 11 10 9 8 5 4 3 2 1
##
## $dev
## [1] 37.33907 37.29157 38.33699 37.65415 38.33797 37.92118 38.40678 37.02964
## [9] 35.68048 33.90672 35.72537 36.70415 35.61927 36.51040
##
## $k
## [1] -Inf 0.4053175 0.4861111 0.6300000 0.7067227 0.7401389 0.8648868
## [8] 0.9463416 1.0285714 1.0544949 1.7640000 2.1310744 2.8578339 3.4374900
##
## $method
## [1] "deviance"
##
## attr("class")
## [1] "prune" "tree.sequence"


```

Plots were developed in order to see the results from performing cross-validation on this pruned tree.


```

plot(Pruned_Netflix$size, Pruned_Netflix$dev, type = "b",
     xlab = "Tree Size", ylab = "SSE")

```



```
plot(Pruned_Netflix$k, Pruned_Netflix$dev, type = "b",  
     xlab = "Cost-Complexity Parameter", ylab = "SSE")
```



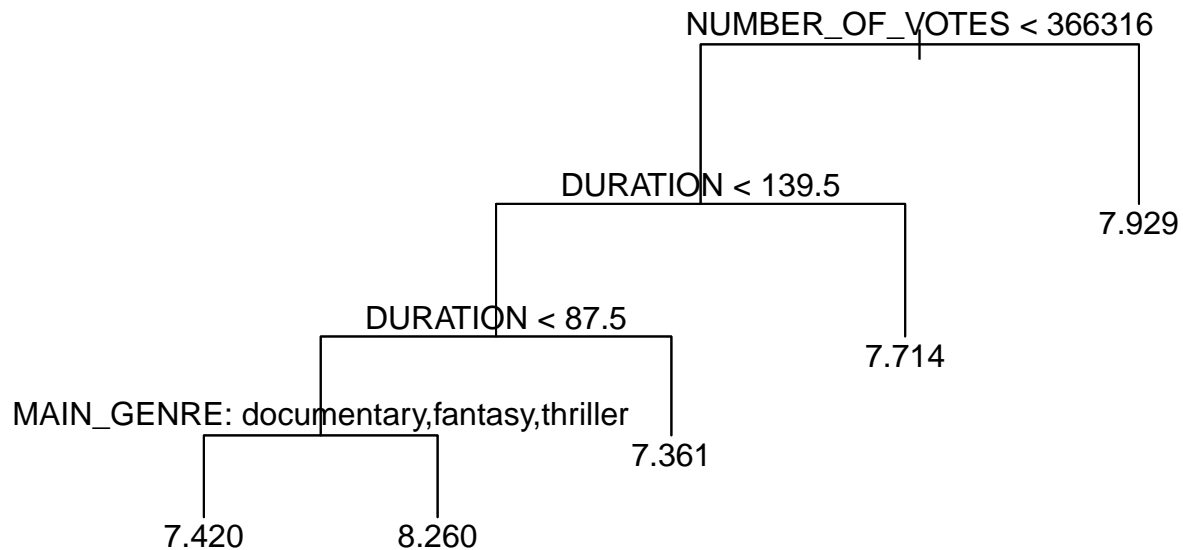
We should choose the tree with the lowest error, which is the tree with 5 nodes. Then, we will predict on the test dataset.

```
NetflixLowestErrorTree = prune.tree(Netflix_regressiontree, best = 5)
summary(NetflixLowestErrorTree)
```

```
##
## Regression tree:
## snip.tree(tree = Netflix_regressiontree, nodes = c(3L, 5L, 9L
## ))
## Variables actually used in tree construction:
## [1] "NUMBER_OF_VOTES" "DURATION" "MAIN_GENRE"
## Number of terminal nodes: 5
## Residual mean deviance: 0.1309 = 24.61 / 188
## Distribution of residuals:
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## -0.72940 -0.26110 -0.01429  0.00000  0.23890  0.98570
```

We will create a plot of this pruned tree.

```
plot(NetflixLowestErrorTree)
text(NetflixLowestErrorTree, pretty = 0)
```



This tree will be used to form the predictions on the test data set.

```

predictedbestNextflixtree <-predict(NetflixForest, newdata=NetflixForest_test)
predictedbestNextflixtree

```

##	1	2	3	4	5	6	7	8
##	7.420000	7.929412	7.929412	7.929412	8.260000	7.714286	7.361069	7.714286
##	9	10	11	12	13	14	15	16
##	7.714286	8.260000	7.929412	7.714286	7.361069	7.361069	7.714286	7.361069
##	17	18	19	20	21	22	23	24
##	8.260000	7.361069	7.929412	7.714286	7.714286	7.420000	7.929412	8.260000
##	25	26	27	28	29	30	31	32
##	7.714286	7.714286	7.714286	7.361069	7.361069	7.361069	7.420000	7.714286
##	33	34	35	36	37	38	39	40
##	7.929412	7.714286	7.929412	7.361069	7.929412	7.714286	7.840000	7.361069
##	41	42	43	44	45	46	47	48
##	7.929412	7.361069	7.929412	7.929412	7.420000	8.260000	7.361069	8.260000
##	49	50	51	52	53	54	55	56
##	7.714286	7.714286	7.929412	7.929412	7.929412	7.361069	7.929412	7.929412
##	57	58	59	60	61	62	63	64
##	7.361069	7.361069	7.714286	7.361069	7.361069	7.361069	7.361069	7.361069
##	65	66	67	68	69	70	71	72
##	7.361069	7.420000	7.929412	7.714286	7.361069	7.929412	7.361069	7.361069
##	73	74	75	76	77	78	79	80
##	7.361069	7.714286	7.361069	7.714286	7.361069	7.361069	7.361069	7.840000

```
##      81      82      83      84      85      86      87      88
## 7.929412 7.361069 7.361069 7.361069 7.714286 7.714286 7.361069 7.361069
##      89      90      91      92      93      94      95      96
## 7.714286 7.361069 7.361069 7.361069 7.929412 7.361069 7.361069 7.361069
##      97      98      99     100     101     102     103     104
## 7.361069 7.361069 7.714286 7.714286 7.714286 7.361069 7.361069 7.361069
##     105     106     107     108     109     110     111     112
## 7.361069 7.929412 7.361069 7.361069 7.714286 7.361069 7.714286 7.361069
##     113     114     115     116     117     118     119     120
## 7.361069 7.361069 7.714286 7.361069 7.420000 7.361069 7.714286 7.361069
##     121     122     123     124     125     126     127     128
## 7.361069 7.361069 7.714286 7.361069 7.714286 7.361069 7.840000 7.361069
##     129     130     131     132     133     134     135     136
## 7.714286 7.714286 7.361069 7.420000 7.361069 7.361069 7.361069 7.361069
##     137     138     139     140     141     142     143     144
## 7.361069 7.361069 7.361069 7.840000 7.361069 7.361069 7.361069 7.714286
##     145     146     147     148     149     150     151     152
## 7.361069 7.361069 7.361069 7.361069 7.361069 7.361069 7.361069 7.361069
##     153     154     155     156     157     158     159     160
## 7.361069 7.361069 7.714286 8.260000 7.361069 7.361069 7.361069 7.714286
##     161     162     163     164     165     166     167     168
## 7.361069 7.714286 7.714286 7.361069 7.361069 7.361069 7.361069 7.361069
##     169     170     171     172     173     174     175     176
## 7.361069 7.361069 7.361069 7.361069 7.361069 7.361069 7.714286 7.361069
##     177     178     179     180     181     182     183     184
## 7.714286 7.361069 7.714286 7.361069 7.714286 7.361069 7.361069 7.714286
##     185     186     187     188     189     190     191     192
## 7.361069 7.361069 7.361069 7.714286 7.361069 7.361069 7.714286 7.361069
##     193     194
## 7.361069 7.361069
```

We will compute the RMSE through the creation of this function.

```
rmse<-function(actual, predicted){
  rmse=sqrt(mean((actual - predicted) ^ 2))
  mse= mean((actual-predicted)^2)
  c(rmse,mse)
}
```

The performance of this tree will be evaluated through using RMSE.

```
rmse(NetflixFTest_set$SCORE,predictedbestNextflixtree)
```

```
## [1] 0.4234998 0.1793521
```

We will load this library in order to perform random forest and bagging.

```
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'

## The following object is masked from 'package:dplyr':
##
##      combine

## The following object is masked from 'package:ggplot2':
##
##      margin
```

We will perform bagging with atleast 500 trees.

```
set.seed(458)

bagging_Netflix <- randomForest(SCORE ~ ., data = NetflixTrain_set, mtry = 4, importance = TRUE, ntree = 500)

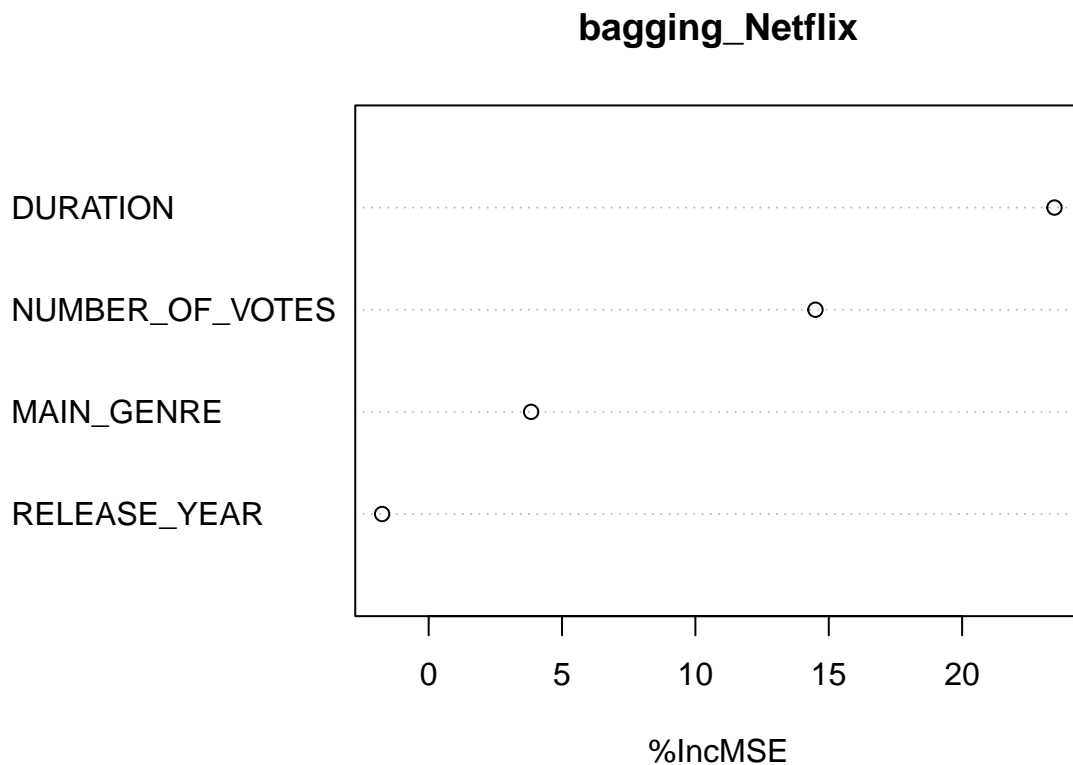
bagging_Netflix
```

```
##
## Call:
## randomForest(formula = SCORE ~ ., data = NetflixTrain_set, mtry = 4, importance = TRUE, ntree = 500)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 0.1647302
##              % Var explained: 8.64
```

```
importance(bagging_Netflix, type = 1)
```

```
##              %IncMSE
## RELEASE_YEAR    -1.744600
## NUMBER_OF_VOTES 14.500684
## DURATION        23.462023
## MAIN_GENRE       3.841185
```

```
varImpPlot(bagging_Netflix, type = 1)
```

The test data set will be used for the predictions.

```
baggingpredictions_Netflix <- predict(bagging_Netflix,newdat=NetflixTest_set)
rmse(NetflixTest_set$SCORE, baggingpredictions_Netflix)
```

```
## [1] 0.4102320 0.1682903
```

Random forest will now be implemented on the Netflix dataset. Since this is regression, the total number of predictors divided by 3 will be the value that is selected for the mtry function.

```
set.seed(458)

Netflix_randomforests <- randomForest(SCORE~.,data=NetflixTrain_set,mtry=1.333333,importance=TRUE,ntree=500)
Netflix_randomforests

##
## Call:
## randomForest(formula = SCORE ~ ., data = NetflixTrain_set, mtry = 1.333333, importance = TRUE,
##               Type of random forest: regression
##               Number of trees: 500
## No. of variables tried at each split: 1
##
##               Mean of squared residuals: 0.1632537
##               % Var explained: 9.46
```

The predictions will be developed on the test set now.

```
randomforestpredictions_Netflix <-predict(Netflix_randomforests,newdat=NetflixTest_set)
rmse(NetflixTest_set$SCORE, randomforestpredictions_Netflix)
```

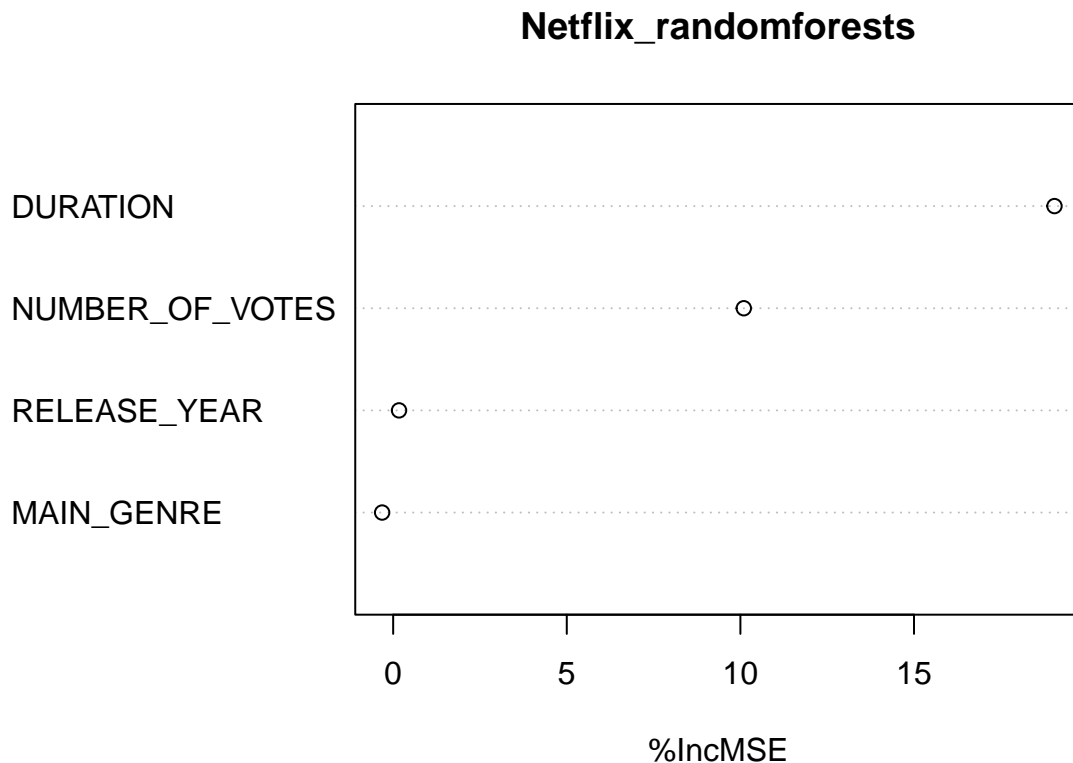
```
## [1] 0.4120424 0.1697789
```

We will begin to examine the importance of each variable and how they operate in the splits of the 500 trees through these two visualizations.

```
importance(Netflix_randomforests,type = 1)
```

```
##              %IncMSE
## RELEASE_YEAR    0.1692325
## NUMBER_OF_VOTES 10.0981472
## DURATION        19.0453766
## MAIN_GENRE      -0.3175460
```

```
varImpPlot(Netflix_randomforests,type = 1)
```



This visual uses the information from random forest's variable importance plot to create a colorful visualization in the form of a bar plot.

```
library(ggplot2)

RF <- data.frame(match = c("Duration", "Number of Votes", "Release Year", "Main Genre"),
                 runs = c(19.0453766, 10.0981472, 0.1692325, -0.3175460))

RandomForestBarPlot <- ggplot(data = RF, aes(x = match, y = runs, fill = match)) +
  geom_bar(stat = "identity", width = 0.5) +
  labs(y = "%IncMSE",
       x = "Predictors",
       title = "Variable Importance: Random Forest")

RandomForestBarPlot + theme(legend.position = "none")
```

